

DTU



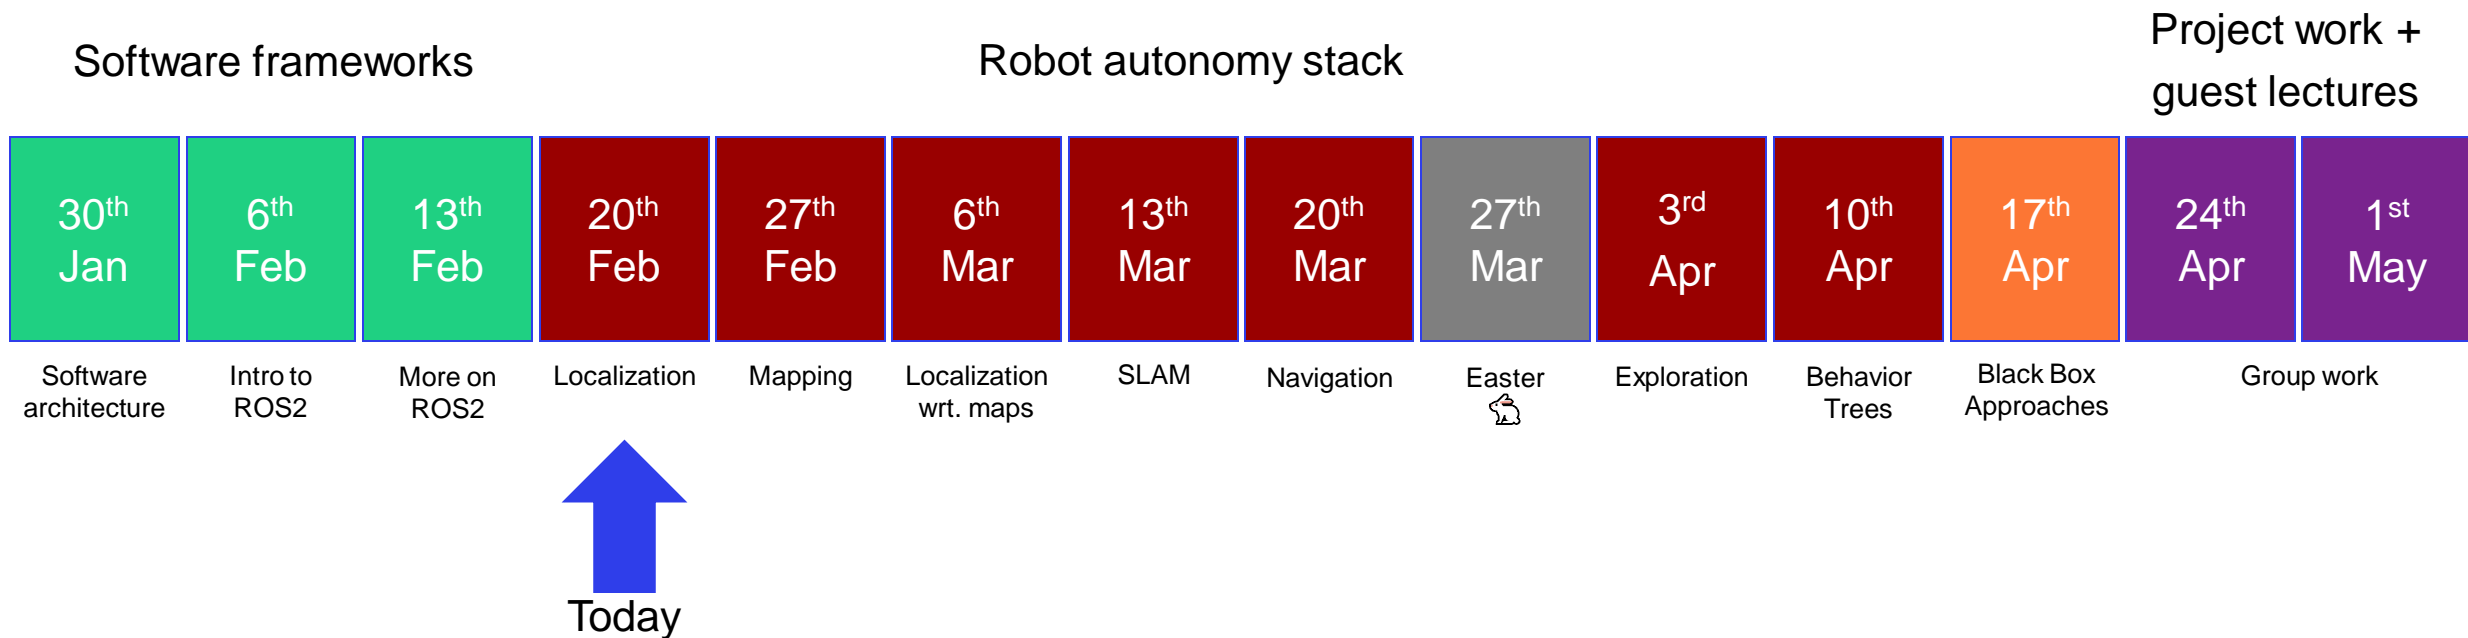
Rasmus Andersen

34761 – Robot Autonomy

(Robot) Localization

Overview of 34761 – Robot Autonomy

- 3 lectures on software frameworks
- 7 lectures on building your own autonomy stack for a mobile robot
- 1 lecture on DL/RL – an overview of black-box approaches to what you have done
- 2 lectures of project work before hand in + guest lectures




Outline for the next 7 weeks

- Our own autonomy stack:

Topic of today

1. Localization
2. Mapping
3. Navigation
4. Exploration
5. Behaviour trees



Outline for the next 7 weeks

- Our own autonomy stack:

1. Localization

- Wheel Odometry
- Visual Odometry
- Visual Inertial Odometry
- Iterative closest point
- LIDAR Odometry


2. Mapping

3. Navigation

4. Exploration

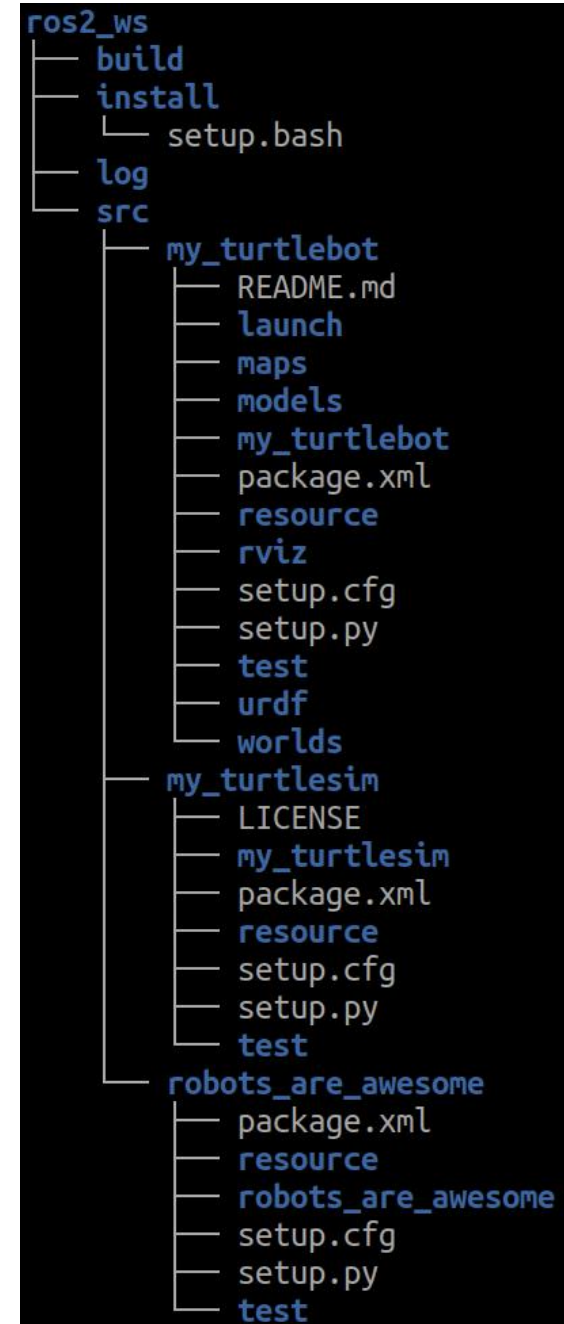
5. Behaviour trees

Topic of today




Recall from last lecture

- A workspace consists of a **build**, **install**, **log**, and **src** folder
 - **The build directory**
 - Intermediate/temporary build files
 - **The install directory**
 - Where the final binaries, libraries, resources, etc. are stored
 - Never modify anything in this directory – it is intended to only be modified by the build system
 - **The log directory**
 - Default location for ROS2 logs
 - **The src directory**
 - Where we put everything to create nodes
 - The only directory we should modify anything in



Recall from last lecture

- Creating ROS nodes to control turtles
 - Anyone up for showing their solution?
- Launching the turtlebot simulation environment (for Mac users with ARM processors)
 - `docker run -p 6080:80 --security-opt seccomp=unconfined --shm-size=512m --platform linux/amd64 tiryoh/ros2-desktop-vnc:humble`
 - Lecture 1 “Before the lecture” section has been updated
 - Your environment should rebuild
 - Be sure to backup your progress and copy it over to the new container
- Transforms





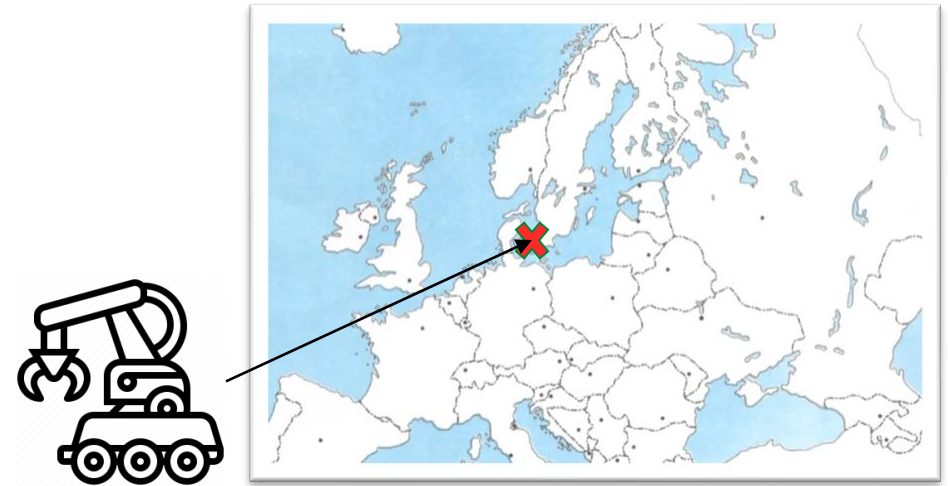
Three students, scripts awry, forgot to heed,
Setup.bash neglected, a crucial feed.
Before the class, they stand with dread,
Singing Rihanna's "Work" in hues of red.



Karaoke 🎵 Work ft. Drake - Rihanna 【No Guide Melody】 Instrumental, Lyric, BGM

What does it mean to localize

- It depends..
 - What do we want to know the location of?
 - Ourselves
 - Our tool
 - An object
 - Localize with respect to what?
 - Do we need global or relative localization?
 - What kind of robot do we have?
 - What kind of sensors do we have?




Global localization

- Just use a GPS




Global localization

- Just use a GPS?
 - What about orientation then?
 - Accuracy?
 - What if this is our robot setup?
 - Or if this is our environment?




- What does global localization even mean?
 - We usually call it **ABSOLUTE** localization, because it's with respect to some defined world reference frame



Forward kinematics of a robot arm

- Use sensor measurement to determine the location of the TPC
- Standard construction for robot manipulators: Denavit-Hartenberg
 - The location is given by a series of matrix multiplications

$${}^{i-1}T_i = \begin{bmatrix} \cos \theta_i & -\sin \theta_i \cos \alpha_{i,i+1} & \sin \theta_i \sin \alpha_{i,i+1} & a_{i,i+1} \cos \theta_i \\ \sin \theta_i & \cos \theta_i \cos \alpha_{i,i+1} & -\cos \theta_i \sin \alpha_{i,i+1} & a_{i,i+1} \sin \theta_i \\ 0 & \sin \alpha_{i,i+1} & \cos \alpha_{i,i+1} & d_i \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



- d_1 Distance between ground and Joint 1
- a_1 Angle of rotation around Joint 1
- d_2 Distance between Joint 1 and Joint 2
- a_2 Angle of rotation around Joint 2
- d_e Distance between Joint 2 and the end effector

Forward kinematics of a robot arm

- Use sensor measurement to determine the location of the TCP
- Standard construction for robot manipulators: Denavit-Hartenberg
 - The location is given by a series of matrix multiplications



Forward kinematics of a robot arm

- Use sensor measurement to determine the location of the TCP
- Standard construction for robot manipulators: Denavit-Hartenberg
 - The location is given by a series of matrix multiplications
- What are the assumptions for this to work?
 - Static environment
 - Objects have the exact same location every time



How does this transfer to mobile robots?


- Use sensor measurement to determine the location of the ~~TPC~~ mobile robots
- The location is given by a series of matrix multiplications

$${}^{i-1}T_i = \begin{bmatrix} \cos \theta_i & -\sin \theta_i \cos \alpha_{i,i+1} & \sin \theta_i \sin \alpha_{i,i+1} & a_{i,i+1} \cos \theta_i \\ \sin \theta_i & \cos \theta_i \cos \alpha_{i,i+1} & -\cos \theta_i \sin \alpha_{i,i+1} & a_{i,i+1} \sin \theta_i \\ 0 & \sin \alpha_{i,i+1} & \cos \alpha_{i,i+1} & d_i \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- So, we just need to keep track of how much we have moved relatively

Relative localization for mobile robots

- Assume no global information about the environment (no map)
- Keep track of movement by integrating over the sensor inputs
 - This is called odometry
- What would the equivalent of measuring joint sensors be for a mobile robot?



Wheel Odometry

- Use wheel encoders to measure how many revolutions a wheel has done
- The size of the wheel determines the distance driven per revolution
- For a differential drive robot:
 - The current position of the robot $\mathbf{p} = [x, y, \theta]^T$

$$\Delta x = \Delta s \cos\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta y = \Delta s \sin\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta\theta = \frac{\Delta s_r + \Delta s_l}{b}$$

$$\Delta s = \frac{\Delta s_r + \Delta s_l}{2}$$

Where b is the distance between the two wheels



Wheel Odometry

- For a differential drive robot:
 - The current position of the robot $\mathbf{p} = [x, y, \theta]^T$

$$\Delta x = \Delta s \cos\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta y = \Delta s \sin\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta\theta = \frac{\Delta s_r + \Delta s_l}{b}$$

$$\Delta s = \frac{\Delta s_r + \Delta s_l}{2}$$

Where b is the distance between the two wheels



Wheel Odometry

- For a differential drive robot:
 - The current position of the robot $\mathbf{p} = [x, y, \theta]^T$

$$\Delta x = \Delta s \cos\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta y = \Delta s \sin\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta\theta = \frac{\Delta s_r + \Delta s_l}{b}$$

$$\Delta s = \frac{\Delta s_r + \Delta s_l}{2}$$

Where b is the distance between the two wheels

$$\mathbf{p}' = \mathbf{p} + \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta\theta \end{bmatrix}$$

By using the relationship between Δs and $\Delta\theta$, we can substitute this further



Wheel Odometry

- For a differential drive robot:

- The current position of the robot $\mathbf{p} = [x, y, \theta]^T$

$$\Delta x = \Delta s \cos\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta y = \Delta s \sin\left(\theta + \frac{\Delta\theta}{2}\right)$$

$$\Delta\theta = \frac{\Delta s_r - \Delta s_l}{b}$$

$$\Delta s = \frac{\Delta s_r + \Delta s_l}{2}$$

Where b is the distance between the two wheels

By using the relationship between Δs and $\Delta\theta$, we can substitute this further:

$$\mathbf{p}' = \mathbf{p} + \begin{bmatrix} \frac{\Delta s_r + \Delta s_l}{2} \cos\left(\theta + \frac{\Delta s_r - \Delta s_l}{2b}\right) \\ \frac{\Delta s_r + \Delta s_l}{2} \sin\left(\theta + \frac{\Delta s_r - \Delta s_l}{2b}\right) \\ \frac{\Delta s_r - \Delta s_l}{b} \end{bmatrix}$$

Wheel Odometry – errors and noise

- Wheel odometry is extremely sensitive to noise
 - Slipping wheels
 - Uneven terrain
 - Misalignment of wheels
 - Shifting contact point of the wheel
 - Error propagation
- In practice, we don't use this method a lot when other methods are available
 - If used, it's often combined with other approaches

Visual Odometry

- Utilize visual features to track changes in localization
- Camera independent – Can be used with a large range of sensors
 - Though, depending on the type of camera, we might have to do some extra steps
- Requires a sequence of images from a rigidly mounted camera
 - Match features in two subsequent images to generate the transformation

- We are specifically looking for the transformation between two subsequent frames:

$$- T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix}$$

- The relative motion between a set of subsequent transformations is then
 - $T_{0:N} = \{T_0, \dots, T_N\}$

- We can then get the current pose through simple matrix multiplications
 - $C_N = C_{N-1}T_N$

Visual Odometry

- We are specifically looking for the transformation between two subsequent frames:

$$- T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix}$$

- The relative motion between a set of subsequent transformations is then


$$- T_{0:N} = \{T_0, \dots, T_N\}$$

- We can then get the current pose simply through simple matrix multiplications

$$- C_N = C_{N-1}T_N$$



- Producing this relative transformation is the core of VO

- Which means we recover the path incrementally
- (though we might perform bundle adjustment as a refinement step)



So how do we produce this relative transform?

- Feature correspondence
 - Use a feature detector on both images and generate correspondence
 - The features must be robust for best performance!
- Appearance based approach
 - Match pixel intensities instead of features
 - Computationally heavier and generally performs worse
 - Not used a lot





So how do we produce this relative transform?

- There are three options for feature-based approaches:
 - 2D to 2D
 - All features are produced in 2D image coordinates
 - 3D to 2D
 - Features of one image plane are given in 3D coordinates and projected to the second image plane
 - 3D to 3D
 - All features are given in 3D – this requires triangulation, e.g., with a stereo camera

The trigonometry

- Assuming a pinhole-model
 - i.e. the mapping from world coordinates to image coordinates is a linear projection
 - $[X, Y, Z] \rightarrow \left[\frac{fX}{Z}, \frac{fY}{Z} \right]$




2D to 2D

- Find the transformation T_k that minimizes the reprojection error:

$$T_k = \operatorname{argmin}_{X_i, C_k} ||p_k - g(X_i, C_k)||^2$$

- Where p_k are our features from image k and $g(X_i, C_k)$ is the reprojection of our features into image plane $k - 1$
- Use the essential matrix to get the reprojection:
 - A minimum of 5 points needed – more = more better
 - p' in image plane O_1 is $Rp' + t$
 - Since $Rp' + T$ and T are on the same (epipolar) plane, the crossproduct $t \times (Rp' + t)$ is normal to the plane
 - p also lies on the epipolar plane! Its dot product should be zero:

$$p \cdot (t \times (Rp' + t)) = 0$$



$$T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix}$$

2D to 2D

- Use the essential matrix to get the reprojection:
 - p' in image plane O_1 is $Rp' + t$
 - Since $Rp' + t$ and t are on the same (epipolar) plane, the crossproduct $t \times (Rp' + t)$ is normal to the plane
 - p also lies on the epipolar plane! Its dot product should be zero:

$$p \cdot (t \times (Rp' + t)) = 0$$

$$p \cdot (t \times (Rp')) = 0$$

- We can express cross product using matrix multiplication:

$$a \times b = \begin{bmatrix} 0 & -a_z & a_y \\ a_z & 0 & -a_x \\ a_y & a_x & 0 \end{bmatrix} \begin{bmatrix} b_x \\ b_y \\ b_z \end{bmatrix} = [a_{\times}]b$$


- We get

$$p^T \cdot [t_{\times}](Rp') = 0$$

$$p^T [t_{\times}] R p' = 0$$

- The matrix $E = [t_{\times}]R$ is the essential matrix!

$$p^T E p' = 0$$



$$T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix}$$

2D to 2D

- So the matrix $E = [t_x]R$ is the essential matrix, and it can be decomposed to a translation and rotation
- Why isn't $E = T_k$ then?
- Scale ambiguity! – we are only looking at epipolar LINES
 - The Essential matrix is a 3×3 matrix that contains 5 degrees of freedom. It has rank 2 and is singular
 - One solution is to use triangulation with previous frames and compute the scale using

$$r = \frac{\|X_{k-1}^1 - X_{k-1}^2\|}{\|X_k^1 - X_k^2\|}$$

Where X are the triangulated features in 3D

- Do this for many points to increase robustness

2D to 2D



Lazaros Nalpantidis

Evangelos Boukas

and it can be decomposed to a translation

If you found this confusing, don't worry
there's a 10ECTS course on the topic

ar LIN

tains a degree of freedom. It has rank 2

us frames and compute the scale using

$$\frac{\|X_{k-1}\|^2}{\|X_k\|^2}$$



Where X are features in an image plane
– Do this for many points to increase robustness

3D to 2D

- This problem is known as camera resection or PnP (perspective from n points)
- Determine the transformation that minimizes the reprojection error

$$T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix} = \underset{T_k}{\operatorname{argmin}} \left\| p_k - p'_{k_1} \right\|^2$$

- Basically boils down to




3D to 3D

- To do this, we need two stereo cameras
- Match the two pointclouds by minimizing the error

$$T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix} = \underset{T_k}{\operatorname{argmin}} \left\| X_k^1 - T_k X_k^2 \right\|^2$$

Where X are the triangulated features in 3D

- This can be done in many ways
 - RANSAC, ICP, Robust Point matching, kernel correlation, and many more

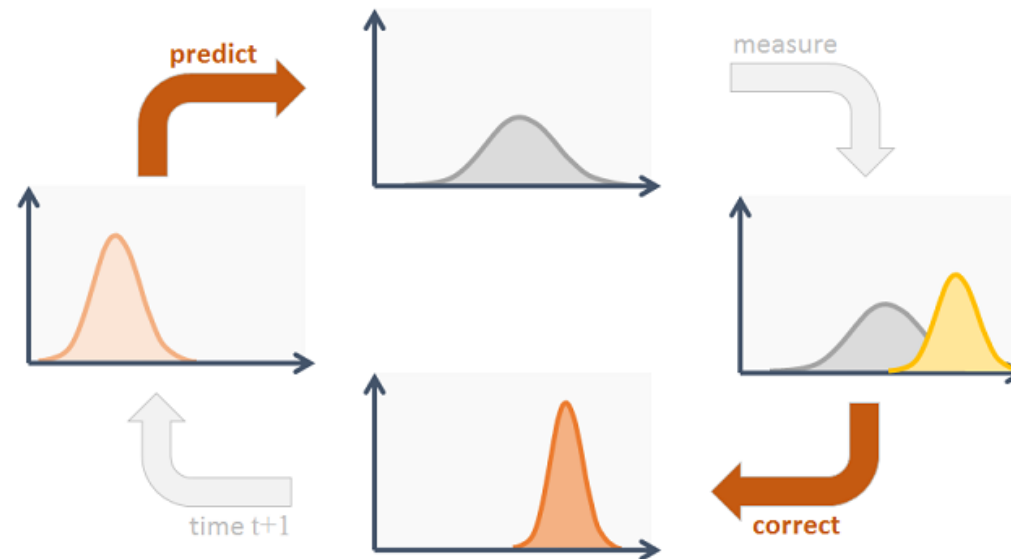


Visual Inertial Odometry

- Visual odometry is great, but...
 - .. we are bound to low-framerates and highly depending on visual features
 - i.e. we can't move fast because of motion blur
 - .. we cannot estimate velocities
- The solution is to fuse our VO with other sensors
 - Inertial Measurement Units provides
 - 6DOF (3DOF acceleration and 3DOF gyroscope) + sometimes a magnetometer
 - Fast sampling rate (can be almost a magnitude faster than the camera FPS)
 - However, they suffer from
 - Sensitive to vibrations
 - Drifts over time – not suitable for localization on its own
- Do you see the combined advantage here?

Visual Inertial Odometry – sensor fusion

- How do we fuse the visual odometry and the IMU?
- We use a Kalman filter
 - Predict the motion from sensor data
 - Update/correct the prediction based on new sensor data



Visual Inertial Odometry – Kalman filter

1. Prediction Step:

- State Prediction: $\hat{x}_{k|k-1} = A\hat{x}_{k-1|k-1} + Bu_k$
- Error Covariance Prediction: $P_{k|k-1} = AP_{k-1|k-1}A^T + Q$

2. Update Step:

- Kalman Gain: $K_k = P_{k|k-1}H^T(H P_{k|k-1}H^T + R)^{-1}$
- State Update: $\hat{x}_{k|k} = \hat{x}_{k|k-1} + K_k(z_k - H\hat{x}_{k|k-1})$
- Error Covariance Update: $P_{k|k} = (I - K_kH)P_{k|k-1}$

Where:

- $\hat{x}_{k|k-1}$ is the predicted state estimate at time k given measurements up to time $k - 1$.
- $\hat{x}_{k|k}$ is the updated state estimate at time k given measurements up to time k .
- A is the state transition matrix.
- B is the control input matrix.
- u_k is the control input at time k .
- $P_{k|k-1}$ is the predicted error covariance matrix at time k given measurements up to time $k - 1$.
- Q is the process noise covariance matrix.
- K_k is the Kalman gain at time k .
- H is the observation matrix.
- R is the measurement noise covariance matrix.
- z_k is the measurement at time k .
- I is the identity matrix.

Iterative closest point


- Matching two pointclouds (like we did for the 3D-3D VO case)

- So we are trying to find T_k

$$T_k = \begin{bmatrix} R_{k-1,k} & t_{k-1,k} \\ 0 & 1 \end{bmatrix} = \underset{T_k}{\operatorname{argmin}} \|X_k - T_k X_k\|^2$$

- This can be solved iteratively for two pointclouds P^1, P^2
 1. For each point p_i^2 , find the nearest point p_j^1
 2. Use all point correspondences to compute T_k^n
 3. Apply T_k^n to P^1
 4. Repeat from step 1 until convergence ($n \rightarrow n + 1$)

Iteration 0



Computing the transformation for ICP


- So how do we actually get the transformation given only the pairs (p_i^1, p_j^2) - Ideas?
- Center the two pointclouds by mean subtraction μ_1, μ_2
- Compute the covariance matrix
 - $C = cov(P'^1, P'^2)$
- Compute the rotation
 - $R = UV^T$ where U and V are SVD components of the covariance matrix
- Compute the translation
 - $t = \mu_1 - R\mu_2$

LIDAR Odometry – KISS ICP

- Keep it small and simple ICP (KISS ICP)
 - Assume a motion/velocity model
 - Can be constant or estimated from previous movement
 - Can also be based on wheel encoders or IMU integration
 - Instead of matching previous LIDAR scan cloud to the current, we try to correct the error from our motion model
 - Use our motion model to predict how much the robot has moved since last scan
 - What are the advantages of doing it this way?
 - We can correct for movements during data acquisition (relevant for fast moving robots like cars)

LIDAR Odometry – KISS ICP

- Keep
 - Ass
 - C
 - C
- Inst
- erro
- U
- Wh
- We
- like



correct the
last scan

moving robots

LIDAR Odometry

- Performing the KISS ICP correction

- Move the pointcloud \hat{P}^* according to your motion model $T_{pred,t}$

$$S = \{s_i = T_{t-1}T_{pred,t} \mathbf{p} | \mathbf{p} \in \hat{P}^*\}$$

- Match the updated scan S with the new pointcloud q

$$\Delta T_{est,j} = \underset{T}{\operatorname{argmin}} ||Ts_i - q||^2$$

$$\{s_i \leftarrow \Delta T_{est,j} s_i | s_i \in S\}$$

- The result is the error of your motion model

$$\Delta T_t = (T_{t-1}T_{pred,t})^{-1} \Delta T_{icp,t} T_{t-1}T_{pred,t}$$

- Note how we applied $T_{pred,t}$ in the local reference frame, we apply $\Delta T_{icp,t}$ in the global reference frame

Exercises

- Create a localization ROS node for your turtlebot
 - Use the LIDAR scanner
 - Assume a constant velocity model
 - Publish a TF with the result of your odometry
 - Compare with the one provided by ROS