

DTU



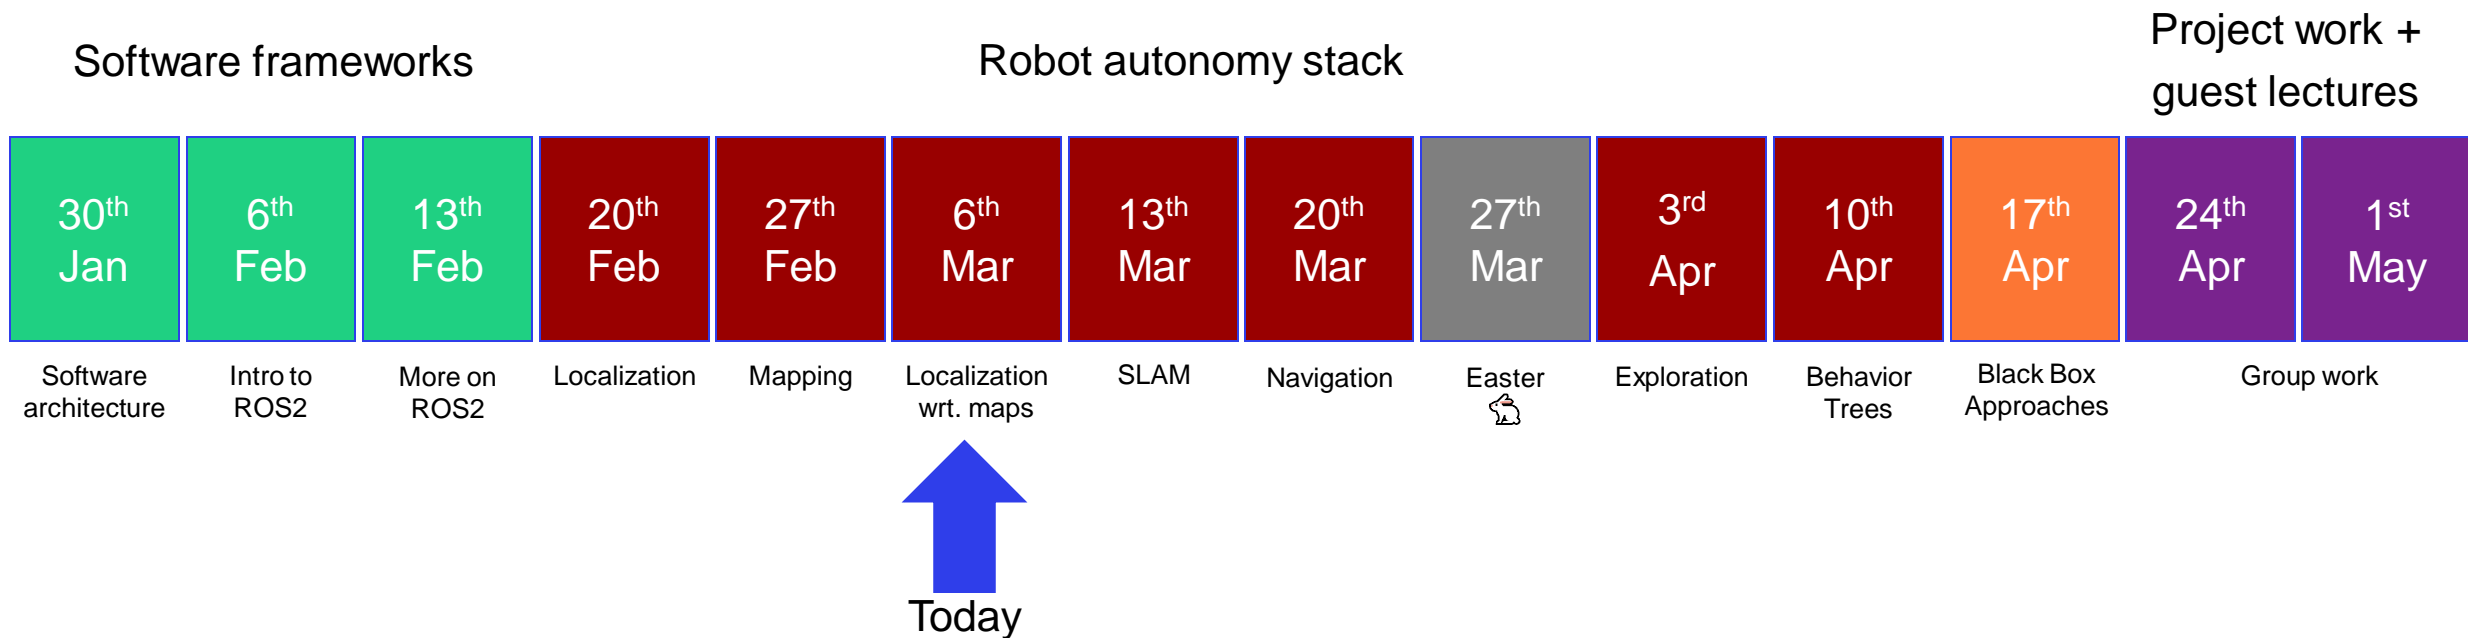
Rasmus Andersen

34761 – Robot Autonomy

Localization w.r.t. maps

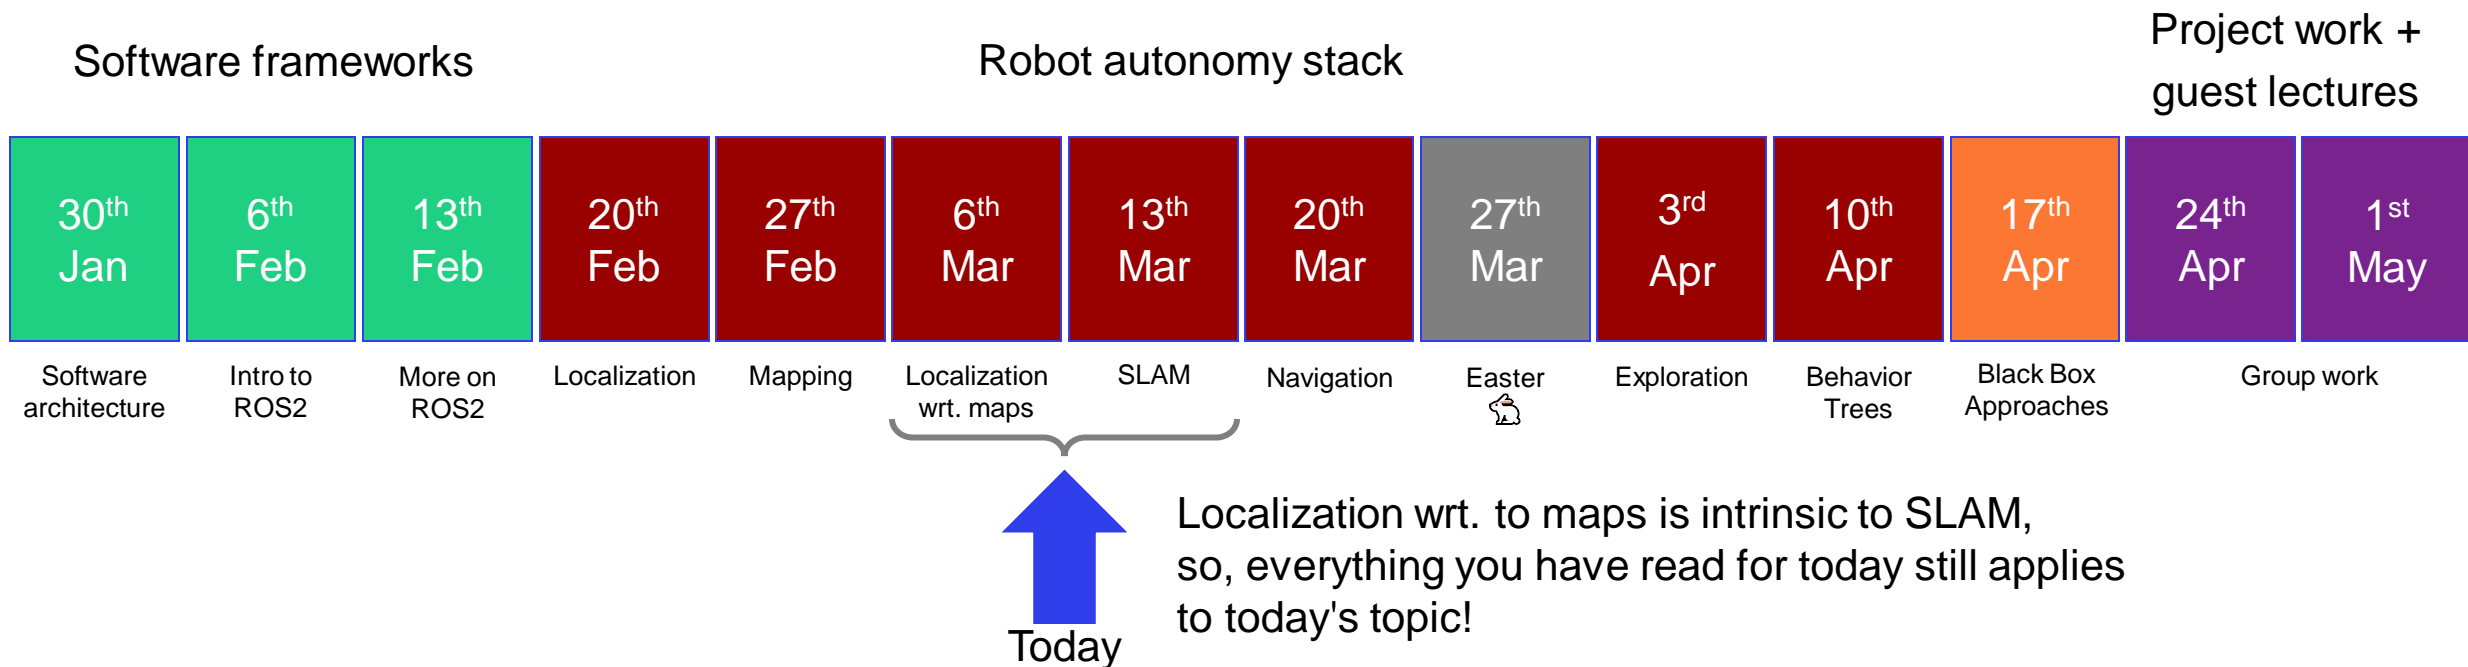
Overview of 34761 – Robot Autonomy

- 3 lectures on software frameworks
- 7 lectures on building your own autonomy stack for a mobile robot
- 1 lecture on DL/RL – an overview of black-box approaches to what you have done
- 2 lectures of project work before hand in + guest lectures



Overview of 34761 – Robot Autonomy

- 3 lectures on software frameworks
- 7 lectures on building your own autonomy stack for a mobile robot
- 1 lecture on DL/RL – an overview of black-box approaches to what you have done
- 2 lectures of project work before hand in + guest lectures




Outline for the next 7 weeks

- Our own autonomy stack:


1. Localization
2. Mapping
3. Navigation
4. Exploration
5. Behaviour trees

Topic of today



Recall from last lecture

- What can we use a map for?
- What are the challenges with a map?
- Map types?
- How can we represent a map?




Exercises from last lecture

- Plotting your sensor data using a ROS interface?
 - Occupied vs free cells
- New forum on LEARN to help answer facilitate questions

Localization w.r.t. a map

- We have a map; we want to know where in the map we are
 - Estimates the location and orientation of the robot in the environment as it moves
- How do we get the initial position?
 - Bayes filtering
 - Particle filter / monte carlo localization



Bayes filtering

- Performs state estimation in a recursive fashion to estimate the current state/location of a system
- From time step t to timestep $t+1$ using only the current observation
- Our belief about the current state

$$Bel(x_t) = P(x_t | z_1, \dots, z_t)$$

- Using bayes rule

$$Bayes = \eta \boxed{P(z_t | x_t, z_1, \dots, z_{t-1})} \boxed{P(x_t | z_1, \dots, z_{t-1})}$$

Likelihood (what's the likelihood of getting z_t)

Prior (our prediction of the state we are in, and therefore, it doesn't depend on z_t)

Bayes filtering

- Using bayes rule


$$\text{Bayes} = \eta P(z_t | x_t, z_1, \dots, z_{t-1}) P(x_t | z_1, \dots, z_{t-1})$$

Likelihood (what's the likelihood of getting z_t)

Prior (our prediction of the state we are in, and therefore, it doesn't depend on z_t)

- Sensor independence
 - Our current sensor input doesn't depend on previous sensor inputs
 - So the likelihood can be simplified to: $P(z_t | x_t)$
 - The prior do depend on previous states, so we can expand the prior:

$$P(x_t | z_1, \dots, z_{t-1}) = \int P(x_t | z_1, \dots, z_{t-1}, x_{t-1}) \cdot P(x_{t-1} | z_1, \dots, z_{t-1}) dx_{t-1}$$



Bayes filtering

- Using bayes rule

$$Bayes = \eta P(z_t | x_t, z_1, \dots, z_{t-1}) P(x_t | z_1, \dots, z_{t-1})$$

Likelihood (what's the likelihood of getting z_t)

Prior (our prediction of the state we are in, and therefore, it doesn't depend on z_t)

- Sensor independence

- Our current sensor input doesn't depend on previous sensor inputs

- So the likelihood can be simplified to: $P(z_t | x_t)$

- The prior do depend on previous states, so we can expand the prior:

$$P(x_t | z_1, \dots, z_{t-1}) = \int P(x_t | z_1, \dots, z_{t-1}, x_{t-1}) \cdot P(x_{t-1} | z_1, \dots, z_{t-1}) dx_{t-1}$$

- The markovian property (i.e. the current state can be explained through only the previous state):

$$\eta P(z_t | x_t) \int P(x_t | z_1, \dots, z_{t-1}, x_{t-1}) \cdot P(x_{t-1} | z_1, \dots, z_{t-1}) dx_{t-1}$$

Effectively, this becomes our environment dynamics

(given x_{t-1} , how likely am I to transition to x_t)

Bayes filtering

- Using bayes rule

$$Bayes = \eta P(z_t | x_t, z_1, \dots, z_{t-1}) P(x_t | z_1, \dots, z_{t-1})$$

Likelihood (what's the likelihood of getting z_t)

Prior (our prediction of the state we are in, and therefore, it doesn't depend on z_t)

- Sensor independence

- Our current sensor input doesn't depend on previous sensor inputs

- So the likelihood can be simplified to: $P(z_t | x_t)$

- The prior do depend on previous states, so we can expand the prior:

$$P(x_t | z_1, \dots, z_{t-1}) = \int P(x_t | z_1, \dots, z_{t-1}, x_{t-1}) \cdot P(x_{t-1} | z_1, \dots, z_{t-1}) dx_{t-1}$$

- The markovian property (i.e. the current state can be explained through only the previous state):


$$\eta P(z_t | x_t) \int P(x_t | z_1, \dots, z_{t-1}, x_{t-1}) \cdot P(x_{t-1} | z_1, \dots, z_{t-1}) dx_{t-1}$$

- $P(x_{t-1} | z_1, \dots, z_{t-1})$ is actually just our belief from the previous timestep: $Bel(x_{t-1})$


$$Bel(x_t) = \eta P(z_t | x_t) \int P(x_t | x_{t-1}) \cdot Bel(x_{t-1}) dx_{t-1}$$

Example



Robot perceives that there are walls above and below, but no walls either left or right




Example




Prob




Example




Prob




Example




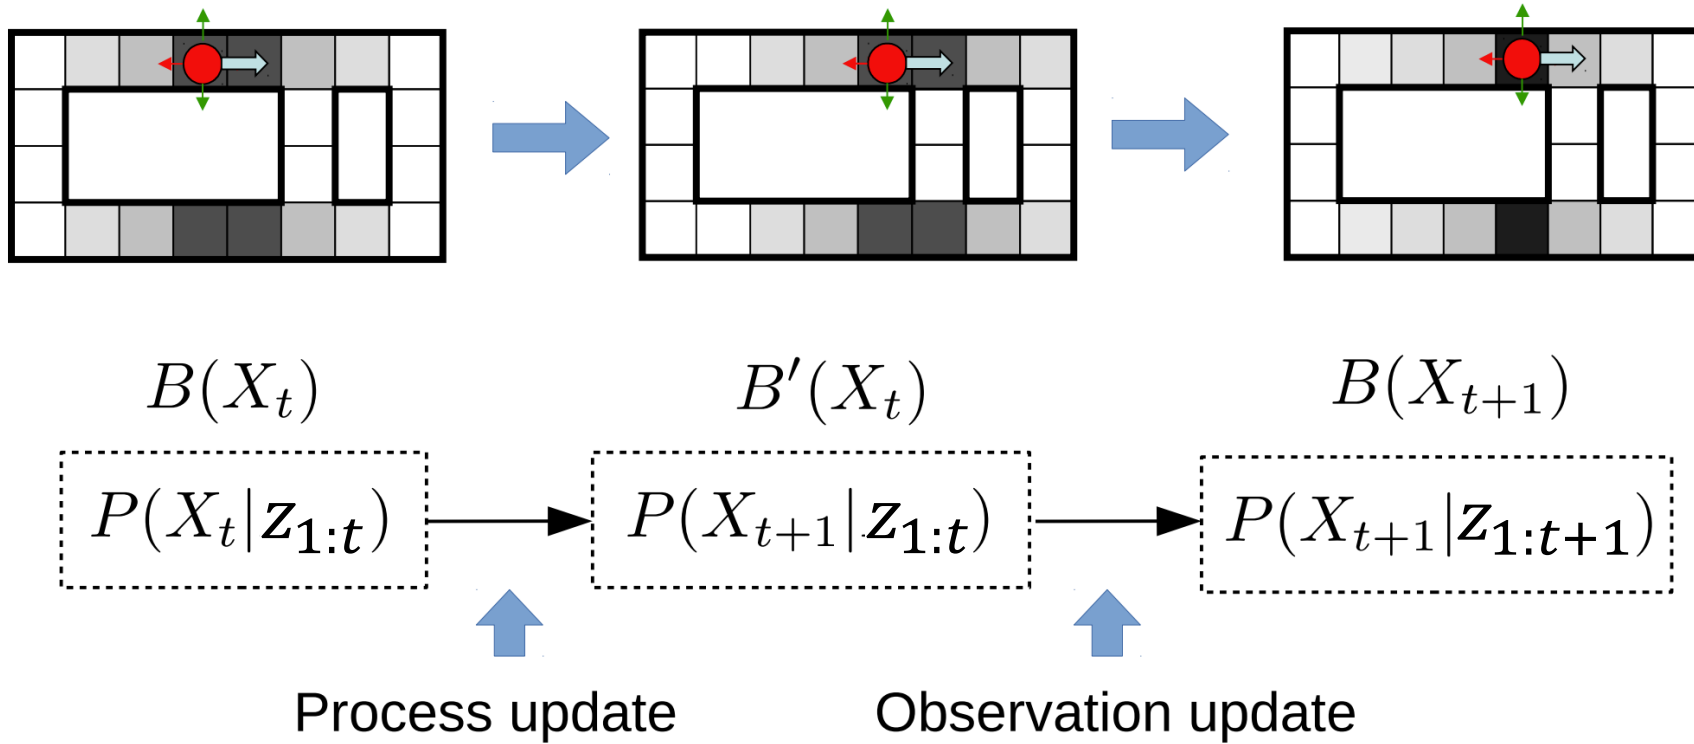
Prob






- The state X_t is unobserved and the one we would like to estimate
- z_t is the observation we can make with our sensors
- From the previous slides we have the prior (observation dynamics) and
$$P(z_t|x_t) \rightarrow \text{Observation dynamics}$$
$$P(x_t|x_{t-1}) \rightarrow \text{Environment dynamics}$$






- Predict the next state
- Correct based on observation


Before process update



Before observation
update





After observation
update




Assumptions we have made

Why must I be confined to this grid?



Particle filter localization

- Represents the location as a distribution of possible states
 - i.e. each particle is a hypothesis of where the robot is
 - Survival of the fittest (particles)
- The initial set of particle hypothesis can be uniformly distributed over the map
- Particle filtering is just an adaptation of the bayes filter using particles instead of grid cells
- Allows us to do non-gaussian distributions as well



Particle filter localization


- Dense particles means higher probability mass
- Weight the importance of each particle to modulate our distribution

- Weighted particles
 - $S = \{(s^i, w^i) | i = 1, \dots, N\}$


State hypothesis
(particle)

Importance weight

- The samples of hypotheses can then be our posterior
 - $P(x) = \sum_{i=1}^N w^i f(s^i)$




Sampling from a distribution




- Rejection sampling
 - Sample x from a uniform distribution
 - Sample c from $[0,1]$
 - If $f(x) > c$ keep the sample, reject otherwise

Sampling from a distribution



- We can even use a different distribution to generate sample from f
 - This is where importance sampling comes in
 - Account for differences in the two distributions:
 - $w = \frac{f}{g}$
 - We upscale the chance of selecting a sample from our proposal distribution according to our target distribution

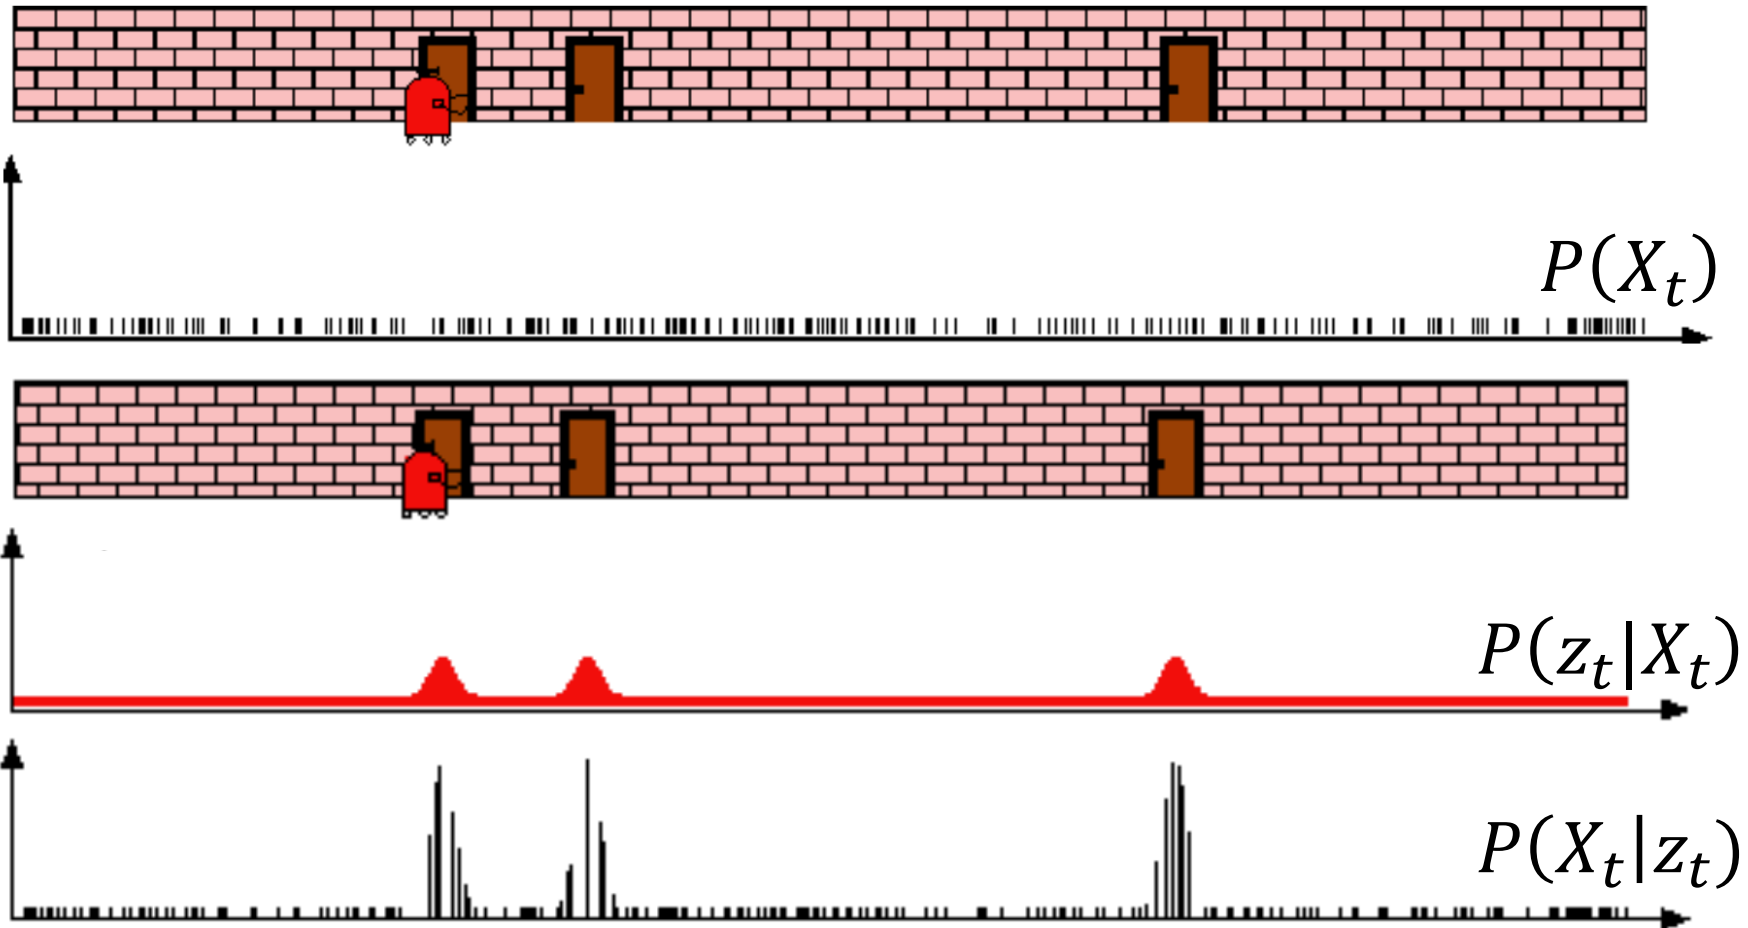
Example



Example

$$Bel(x_{t+1}) = \eta P(z_{t+1}|X_{t+1}) Bel(X_t)$$

$$w = \frac{\eta P(z_{t+1}|X_{t+1}) Bel(X_t)}{Bel(X_{t+1})} = \eta P(z_{t+1}|X_{t+1})$$




The particle filter algorithm

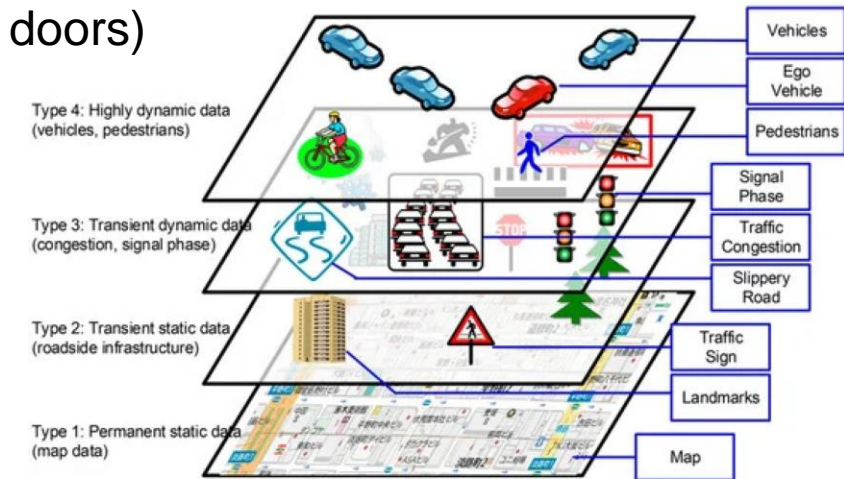
```
Algorithm MCL( $X_{t-1}, u_t, z_t$ ):  
   $\bar{X}_t = X_t = \emptyset$   
  for  $m = 1$  to  $M$ :  
     $x_t^{[m]} = \text{motion\_update}(u_t, x_{t-1}^{[m]})$   
     $w_t^{[m]} = \text{sensor\_update}(z_t, x_t^{[m]})$   
     $\bar{X}_t = \bar{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
  endfor  
  for  $m = 1$  to  $M$ :  
    draw  $x_t^{[m]}$  from  $\bar{X}_t$  with probability  $\propto w_t^{[m]}$   
     $X_t = X_t + x_t^{[m]}$   
  endfor  
  return  $X_t$ 
```

What do we need to be aware of?

- Distinctive features in the environment are necessary
 - To do the sensor update, we need to compute $P(z_t|x_t)$ which require us to quantify the state
 - This is done through features in the environment (like the doors)



- What if the environment is very dynamic?
 - Utilize the mapping abstractions, such as HD map representations
- High sensitivity to the number of points!
- We need to move – i.e. we need to change the state to get updates



Exercises

- Using your own localization and accumulate a map when you drive around in the environment
 - Use a counter to define if a cell is free or occupied
- The simulation we are using already uses a particle filter to perform localization
 - Set an initial starting point of the robot in rviz using the “2D pose estimate” button
 - Create a ROS2 node that:
 - Prints the number of particles used
 - Computes the expected position of the robot based on the particles
 - Change the number of particles when you launch the simulation
- Implement your own particle filter for localization in a map
 - Use the `cmd_topic` to estimate your motion model
 - Subscribe to the LIDAR topic and compute features
 - For each particle compute the error for each feature
 - Update your importance weights based on the error