

DTU



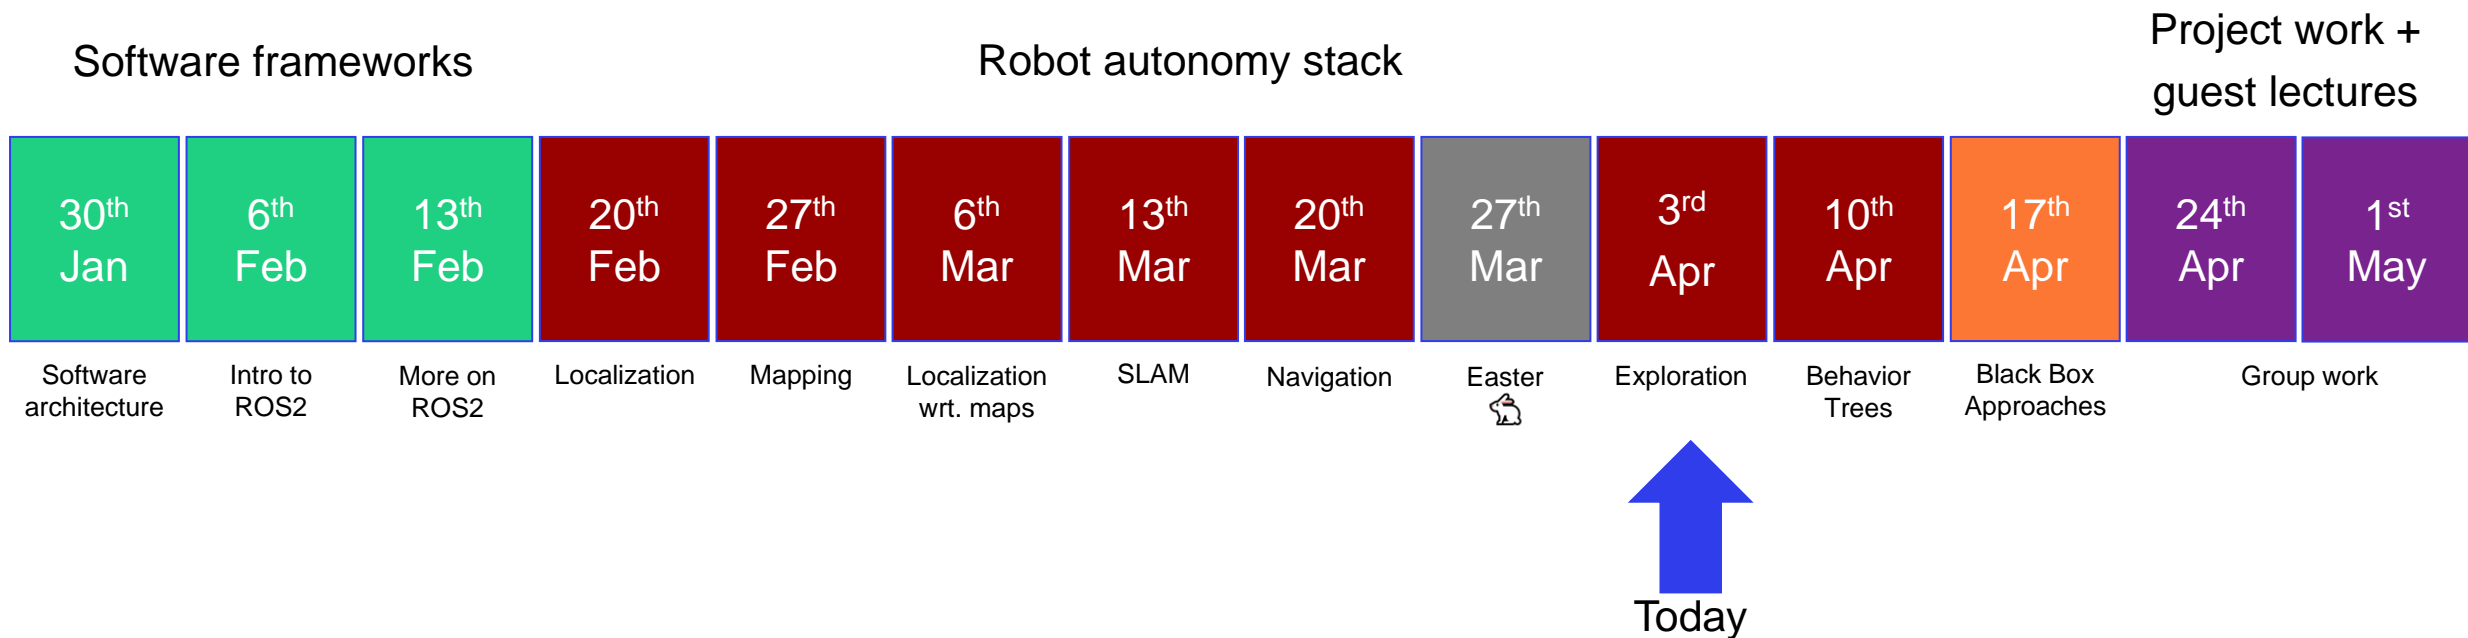
Rasmus Andersen

34761 – Robot Autonomy

# Exploration

# Overview of 34761 – Robot Autonomy

- 3 lectures on software frameworks
- 7 lectures on building your own autonomy stack for a mobile robot
- 1 lecture on DL/RL – an overview of black-box approaches to what you have done
- 2 lectures of project work before hand in + guest lectures




# Outline for the next 7 weeks

- Our own autonomy stack:
  1. Localization
  2. Mapping
  3. Navigation
  4. Exploration
  5. Behaviour trees

Topic of today

4. Exploration

5. Behaviour trees



# Outline for the next 7 weeks

- Our own autonomy stack:


1. Localization
2. Mapping
3. Navigation

4. Exploration

1. Depth-first search
2. Breadth-first search
3. Frontier-based exploration
4. Next-best-view exploration


5. Behaviour trees

Topic of today





# Recall from last lecture – discretization of configuration space


- Configuration spaces are, in general, continuous spaces (e.g.,  $R^2$  in the example)
- We simplify the path planning problem by discretizing the space, e.g.,
  - Gridding
  - Random sampling free configurations ●
- Graphs become powerful computational tools for representing the configuration space



# Graphs



Unweighted undirected graph



Weighted directed graph

# Recall from last lecture – map-based planning algorithms

- Distance transform
- Voronoi roadmap method
- Probabilistic roadmap method
- Dijkstras Algorithm
- Rapidly-exploring random tree (RRT)



# How is exploration different than SLAM?

In SLAM, we try to localize the robot while building a map

- i.e. there is nothing to dictate where or how the robot should move
- Purely passive, and can be done in post-processing

In exploration we try to maximize the map-building autonomously

- Ideally by quantifying what the robot “learns” by navigating to a new pose
- Has to be performed online, the algorithms dictate where the robot navigate to

**NB:** exploration is also different than *coverage*

- Coverage problems assume a map and that we want the robot to optimally cover the map (e.g. vacuum cleaner or lawn mover)

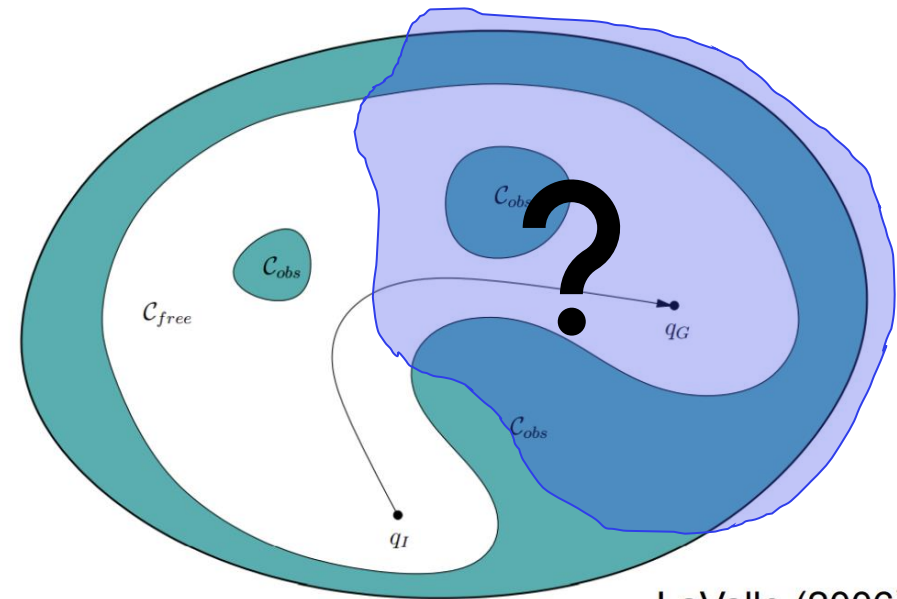
# Exploration applications

- Exploration is central to a range of indoor, outdoor, in-air and underwater applications for autonomous vehicles.
- At home: vacuum cleaner, lawn mower
- Air: surveillance with unmanned air vehicles
- Underwater: reef monitoring
- Underground: exploration of mines
- Space: terrain mapping for localization



# The exploration problem


- We don't have a full map of the environment
  - $\mathcal{C}$  and therefore  $\mathcal{C}_{free}$ ,  $\mathcal{C}_{obs}$  are only partially known
- The goal is to 'maximize'  $\mathcal{C}$ , and less to navigate from point A to point B
  - i.e. produce a map configuration we can use for SLAM and navigation



LaValle (2006)

# The exploration problem

- How do we quantify exploration?
- How do we maximize the exploration?
- Terminology
  - Frontier
    - The border between our open space and unknown space
  - Frontier centroid
    - The center of the frontier border
  - Information gain
    - A metric for how much the robot learns from a given goal pose



# The exploration problem

## Problem definition

The exploration path planning problem consists in exploring a bounded 2D space  $V \subset R^2$ . This is to determine which parts of the initially unmapped space  $V_{unm} = V$  are free  $V_{free} \subset V$  or occupied  $V_{occ} \subset V$ . The operation is subject to vehicle kinematic and dynamic constraints, localization uncertainty and limitations of the employed sensor system with which the space is explored.

- As for most sensors the perception stops at surfaces, hollow spaces or narrow pockets can sometimes not be explored with a given setup. This residual space is denoted as  $V_{res}$ . The problem is considered to be fully solved when  $V_{free} \cup V_{occ} = V/V_{res}$ .
- Due to the nature of the problem, a suitable path has to be computed online and in real-time, as free space to navigate is not known prior to its exploration.

# Exploration approaches for today

- Depth first search & Breadth first search
  - Finding the frontier
  - Exploring the frontier
- Next-best view exploration
  - Quantifying the exploration approach



# Frontier exploration

- The scenario: we have a partial map
- If we can identify the border of where we have been before, we can move towards the unknown



# Frontier exploration




- The scenario: we have a partial map
- If we can identify the border of where we have been before, we can move towards the unknown








# Finding the frontier - Depth first search

- Create a graph of all the cells in your map
- Traverse the branches prioritizing depth
  - NB, beware of loops (non-termination) – keep track of already visited nodes!
  - Sometimes we set a max depth to avoid graphs with loops or infinite depths
- This approach can set high memory requirements due to the tracking of the visited states




# Finding the frontier - Breadth first search

- Explore horizontally instead of vertically in the graph
- Needs to keep track of child nodes that have been queued but not yet visited



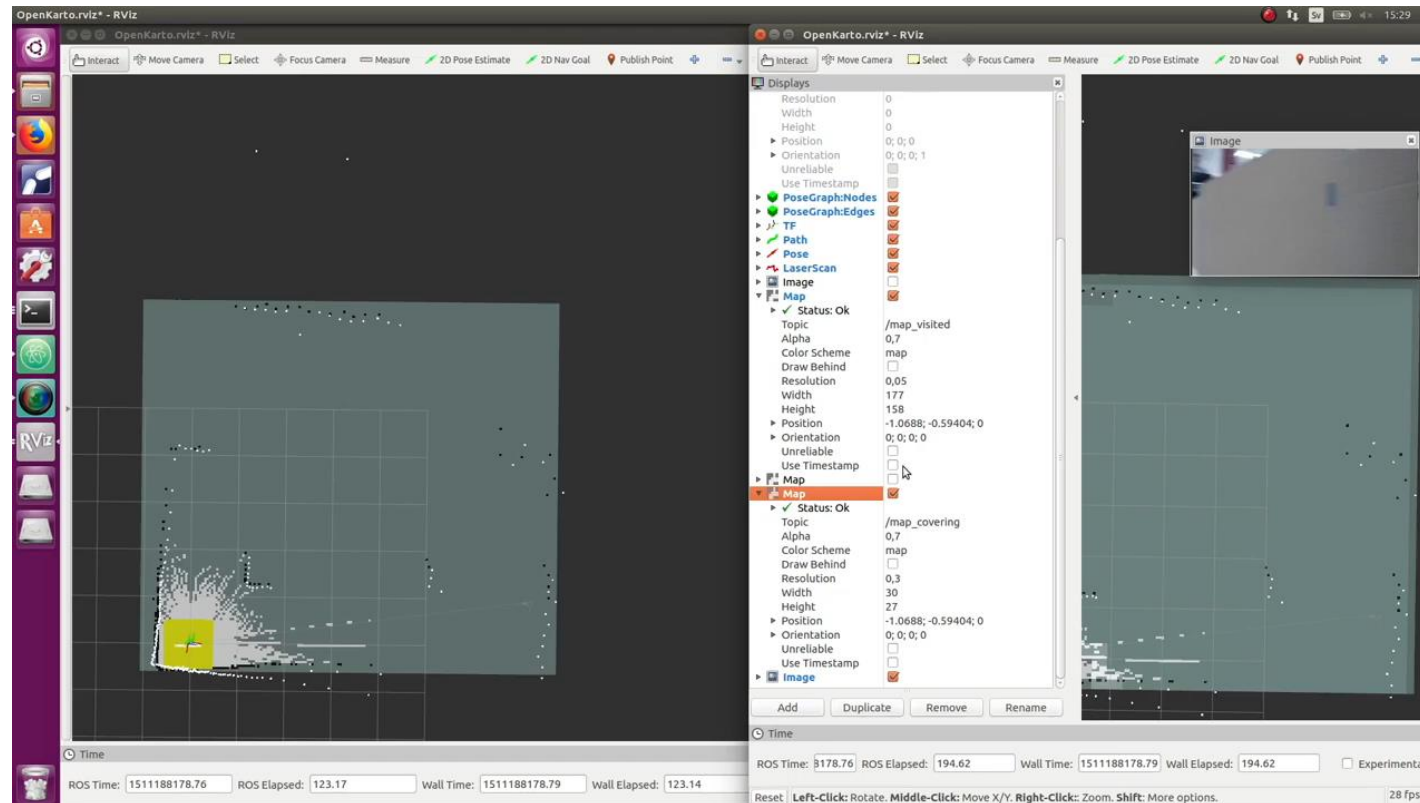
# Depth- and Breadth-first search


- Can be used to find the border between the observed and unknown space
  - i.e. every time we visit an open-space node that is connected to a unknown node, mark it as a frontier
- In a second run of either depth or breadth-first search, only add connected frontier nodes to the graph – i.e. separate frontiers




# Depth- and Breadth-first search

- Can be used to find the border between the observed and unknown space
  - i.e. every time we visit an open-space node that is connected to a unknown node, mark it as a frontier






# Exploring the frontier



- To keep track of multiple frontiers, we perform a second search only on the identified frontier points
- To explore, navigate the robot towards the closest frontier
- Update the map after reaching the frontier, to get new frontiers

# Remarks on exploring the frontier



- A very minimalistic approach
  - See a frontier and move towards it
- Doesn't quantify our goal of exploration
  - Why is this frontier better than all other frontiers?
- Requires us to perform a depth- or breadth-first search after each navigation
  - The map is updated which produces new frontiers than needs to be found
- Information maximization

# Next-best view exploration

- An algorithm for fast exploration of unknown environments
- Defines a sequences of states/nodes from a graph (e.g. build using PRM or RRT)
- Select the sequence of states/nodes with the best sequence of viewpoints
- Execute only the first step of the best sequence (i.e. the path producing the best views of unknown areas)
- Repeat the whole process in a receding horizon fashion until the entire environment has been explored




# Next-best view functional principle

- Tree-based exploration: At every iteration, NBVP spans a random tree of finite depth. Each vertex of the tree is annotated regarding the collected Information Gain – a metric of how much new space is going to be explored.


$$\text{Gain}(n_k) = \text{Gain}(n_{k-1}) + \text{Visible}(M, \xi_k) e^{-\lambda c(\sigma_{k-1}^k)}$$

- Within the sampled tree, evaluation regarding the path that overall leads to the highest information gain is conducted. This corresponds to the best path for the given iteration. It is a sequence of next-best-views as sampled based on the vertices of the spanned random tree.



Kostas Alexis, Robotics Short Seminars, Feb 11 2016


# Next-best view functional principle



$$\mathbf{Gain}(n_k) = \mathbf{Gain}(n_{k-1}) + \mathbf{Visible}(\mathcal{M}, \xi_k) e^{-\lambda c(\sigma_{k-1}^k)}$$


# Next-best view functional principle

- $\text{Gain}(n_k) = \text{Gain}(n_{k-1}) + \text{Visible}(\mathbf{M}, \xi_k) e^{-\lambda c(\sigma_{k-1}^k)}$ 
  - $\text{Gain}(n_{k-1})$  is the information gain for all previous nodes up to the current node
  - $\text{Visible}(\mathbf{M}, \xi_k)$  is the number of cells in the map  $\mathbf{M}$  than can be seen from the robot configuration  $\xi_k$
  - $e^{-\lambda c(\sigma_{k-1}^k)}$  is a discounting factor to limit the horizon of the planner (i.e. receding horizon). This ensure immediate nodes have a bigger influence on the selected path than nodes deep into the node tree
  - $c(\sigma_{k-1}^k)$  is the cost of moving from node  $n_{k-1}$  to  $n_k$  along path  $\sigma_{k-1}^k$  (this could be the Euclidian distance if the path is just a straight line)
  - $\lambda$  is a tuning factor




# Next-best view functional principle

- Environment representation: Occupancy Map dividing space cells that can be marked either as free, occupied or unmapped.
- Only sample robot configurations within the free space our map to avoid planning trajectories into areas we do not know
  - These configurations will generally also have very high gains
- At each viewpoint/configuration of the environment  $\xi$ , the amount of space that is visible is computed as  $\text{Visible}(M, \xi)$ 
  - This means we need efficient ray-casting (e.g. using a camera matrix for depth cameras or projection for LiDARs)




# Next-best view functional principle

- Receding Horizon: For the extracted best path of viewpoints, only the first viewpoint is actually executed.
- The system moves to the first viewpoint of the path of best viewpoints.
- Subsequently, the whole process is repeated within the next iteration. This gives rise to a receding horizon operation.




# Next-best view algorithm

- ♦  $\xi_0 \leftarrow$  current vehicle configuration
- ♦ Initialize  $\mathbf{T}$  with  $\xi_0$  and, unless first planner call, also previous best branch
- ♦  $g_{best} \leftarrow 0$  // Set best gain to zero
- ♦  $n_{best} \leftarrow n_0(\xi_0)$  // Set best node to root
- ♦  $N_T \leftarrow$  Number of nodes in  $\mathbf{T}$
- ♦ **while**  $N_T < N_{max}$  or  $g_{best} == 0$  **do**
  - ♦ Incrementally build  $\mathbf{T}$  by adding  $n_{new}(\xi_{new})$
  - ♦  $N_T \leftarrow N_T + 1$
  - ♦ **if**  $Gain(n_{new}) > g_{best}$  **then**
    - ♦  $n_{best} \leftarrow n_{new}$
    - ♦  $g_{best} \leftarrow Gain(n_{new})$
  - ♦ **if**  $N_T > N_{TOT}$  **then**
    - ♦ Terminate exploration
- ♦  $\sigma \leftarrow \text{ExtractBestPathSegment}(n_{best})$
- ♦ Delete  $\mathbf{T}$
- ♦ **return**  $\sigma$



# Examples of a Next-best view planner





# Example of a Next-best view planner

## Next-Best-View planning for surface reconstruction of large-scale 3D environments with multiple UAVs

Guillaume Hardouin<sup>1,2</sup>, Julien Moras<sup>1</sup>, Fabio Morbidi<sup>2</sup>,  
Julien Marzat<sup>1</sup>, El Mustapha Mouaddib<sup>2</sup>

<sup>1</sup> DTIS, ONERA, Université Paris-Saclay, Palaiseau, France

<sup>2</sup> MIS laboratory, Université de Picardie Jules Verne, Amiens, France



IROS 2020





# Remarks on next-best view planning



- Can be inherently collision free as long as the path  $\sigma_{k-1}^k$  is collision free
  - This means you can use any of the planners from last week (PRM, RRT) and if they produce collision free navigation, the next-best view algorithm will also be collision free
- Most of the computational costs are from the collision checking, however, the information gain computation increases with the resolution of the map
  - Remember, we have to compute the number of visible cells in each robot configuration
- The Receding Horizon Next-Best-View Exploration Planner relies on the real-time update of the 3D map of the environment.
  - We need to update the map, so we can update the gain computation as we explore the map

# Remarks on next-best view planning


- The metric we use doesn't have to be purely unoccupied space – we can add multiple types of information to the information gain

$$Q(n_k) = Q(n_{k-1}) + \mu E(M, \xi) \cdot e^{-\lambda c(\sigma_{k-1}^k)}$$


$$\mu = \sum_{i \in I} \frac{1}{d_i}$$




t=5 min




t=10 min




t=15 min



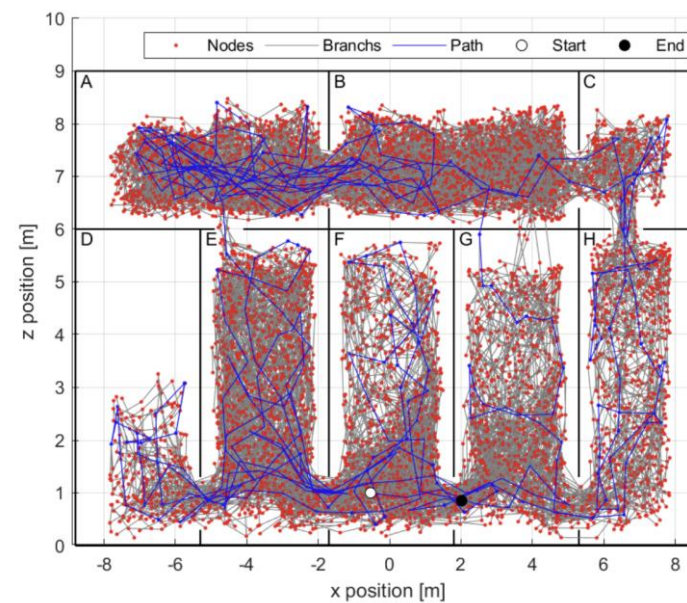
t=20 min



t=25 min



t=30 min

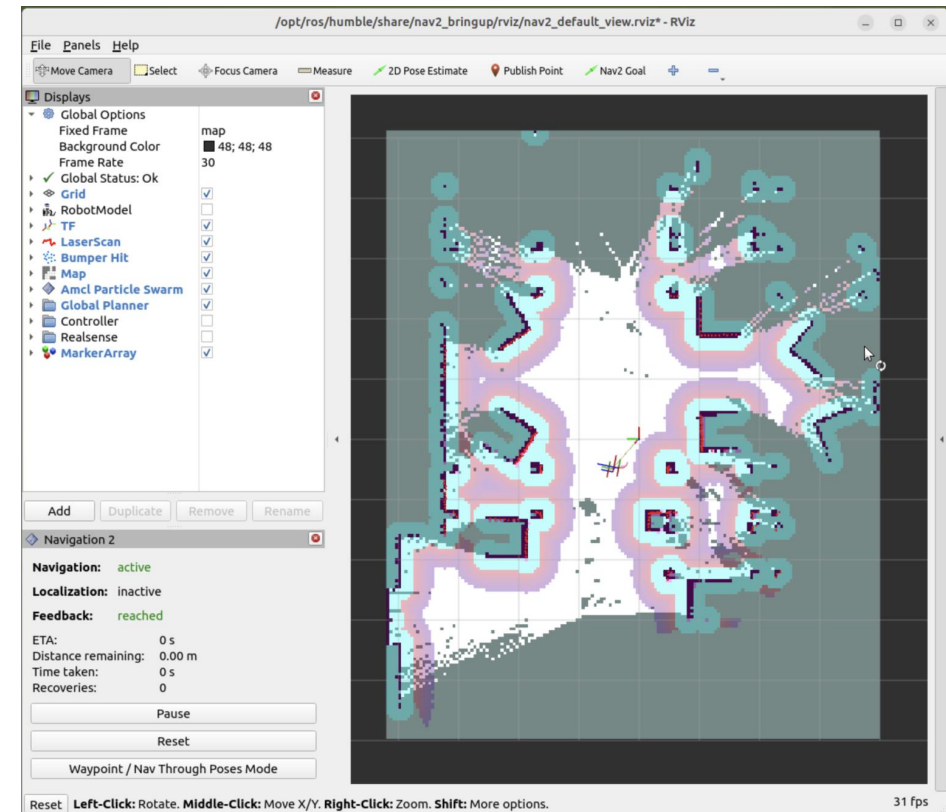


# Exercises

- Create a new map of the environment through rviz, *but only map it partially*
- Build a probabilistic roadmap of the partial map you just created
  - to reduce computational load, you can reject nodes sampled in unknown space
  - Keep it simple initially with just a few random nodes, and visualize them in rviz2
- For each node in the roadmap, compute the information gain according to

$$\text{Gain}(n_k) = \text{Gain}(n_{k-1}) + \text{Visible}(M, \xi_k) e^{-\lambda c(\sigma_{k-1}^k)}$$

- The visible cells has to be computed with ray-casting
- Set  $\lambda = 1$
- You can assume a straight path between the current robot location and the next node and use a Euclidian cost function, i.e.  $c(\sigma_{k-1}^k) = \|\mathbf{x}_k - \mathbf{x}_{k-1}\|_2$



# Exercises

- Execute the path with highest information gain
  - The action interface */navigate\_to\_pose* will drive the robot to the specified location
  - Example using the terminal:

```
ros2 action send_goal /navigate_to_pose
nav2_msgs/action/NavigateToPose '{pose: {header:
{frame_id: "map"}, pose: {position: {x: 0.0, y: 0.0, z:
0.0}, orientation: {x: 0.0, y: 0.0, z: 0.0, w: 1.0}}}}'
```

